



# Adversarial Label Flips

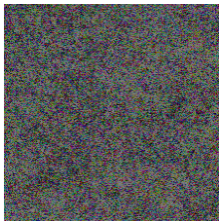
Matthias Dellago & Maximilian Samsinger

A short recap


# Fast gradient sign method



+ €



=



Husky

(42.82% confidence)

Noise (PGD-40)

50x amplified

Handkerchief

(99.999988% confidence)

Source: `ctf.codes`, circa 2021

# What we want to do

## Confusion Matrix

		Categorised as		
		Dog	Cat	Plane
Adversarial Example of a	Dog	0.0	?	?
	Cat	?	0.0	?
	Plane	?	?	0.0

How many modified dogs get classified as cats vs as planes? etc.

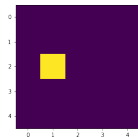
Some simple theory

We want similar images that are classified differently.  
But what is "similar"?

# Quantifying Changes

$L^0$ -Norm

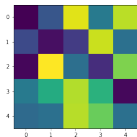
Number of  
pixels changed



Perturb one  
pixel maximally

$L^1$ -Norm

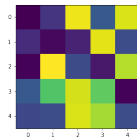
Sum of  
all changes



Minimise sum

$L^2$ -Norm

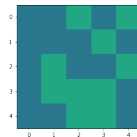
Sum of the *square*  
of all changes



Minimise sum  
of squares

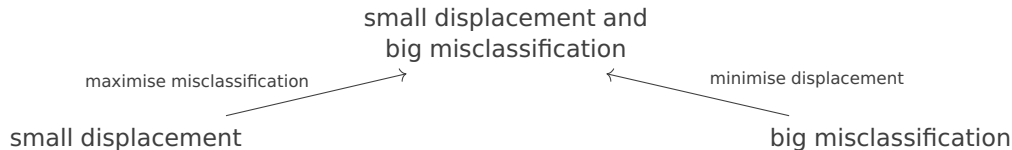
$L^\infty$ -Norm

Maximum of  
all changes



Perturb all  
pixels equally

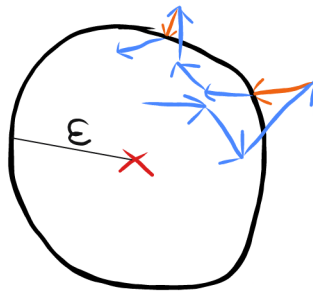
# Two Different Approaches






# Projected Gradient Decent

- 1 Pick spot in epsilon ball
- 2 Iterate gradient decent
- 3 If leaving ball, project back onto surface.



# Projected Gradient Decent



Know your enemy, Oscar Knagg, towardsdatascience.com, 2019

# Carlini-Wagner-Attack

Original idea: minimise distance while always staying in "misclassification territory".

Problem: Nonlinearity of constraint makes for bad optimisation

# Carlini-Wagner-Attack

Solution: Pack constraint into the function that is optimised.  
=> minimise: distance + "how misclassified is x?"\*  
equiv. minimise distance while maximising misclassification.

\*loss function

# References I